

# assignment-3

April 15, 2024

1 Name: Sourabh Barala

2 Course: M.Sc. Data Science

3 Year: 1st

4 Reg. No.: 23MSD7044

5 Subject: Machine Learning and its Applications

6

7

8

9

10

11

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestClassifier, BaggingClassifier, \
    AdaBoostClassifier, GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
```

```

from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
import xgboost as xgb
from sklearn.model_selection import KFold, cross_val_score

```

```
[2]: titanic=pd.read_csv('archive/tested.csv')
```

```
[3]: titanic.head(5)
```

```
[3]:
```

	PassengerId	Survived	Pclass	\
0	892	0	3	
1	893	1	3	
2	894	0	2	
3	895	0	3	
4	896	1	3	

	Name	Sex	Age	SibSp	Parch	\
0	Kelly, Mr. James	male	34.5	0	0	
1	Wilkes, Mrs. James (Ellen Needs)	female	47.0	1	0	
2	Myles, Mr. Thomas Francis	male	62.0	0	0	
3	Wirz, Mr. Albert	male	27.0	0	0	
4	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	22.0	1	1	

	Ticket	Fare	Cabin	Embarked
0	330911	7.8292	NaN	Q
1	363272	7.0000	NaN	S
2	240276	9.6875	NaN	Q
3	315154	8.6625	NaN	S
4	3101298	12.2875	NaN	S

```
[4]: titanic.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  418 non-null    int64
1   Survived     418 non-null    int64
2   Pclass       418 non-null    int64
3   Name         418 non-null    object
4   Sex          418 non-null    object
5   Age         332 non-null    float64
6   SibSp        418 non-null    int64
7   Parch        418 non-null    int64
8   Ticket       418 non-null    object
9   Fare         417 non-null    float64

```

```

10 Cabin          91 non-null    object
11 Embarked       418 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 39.3+ KB

```

```
[5]: titanic.describe()
```

```

[5]:      PassengerId  Survived  Pclass    Age  SibSp  \
count    418.000000   418.000000   418.000000  332.000000  418.000000
mean     1100.500000     0.363636     2.265550   30.272590    0.447368
std       120.810458     0.481622     0.841838   14.181209    0.896760
min        892.000000     0.000000     1.000000    0.170000    0.000000
25%        996.250000     0.000000     1.000000   21.000000    0.000000
50%       1100.500000     0.000000     3.000000   27.000000    0.000000
75%       1204.750000     1.000000     3.000000   39.000000    1.000000
max       1309.000000     1.000000     3.000000   76.000000    8.000000

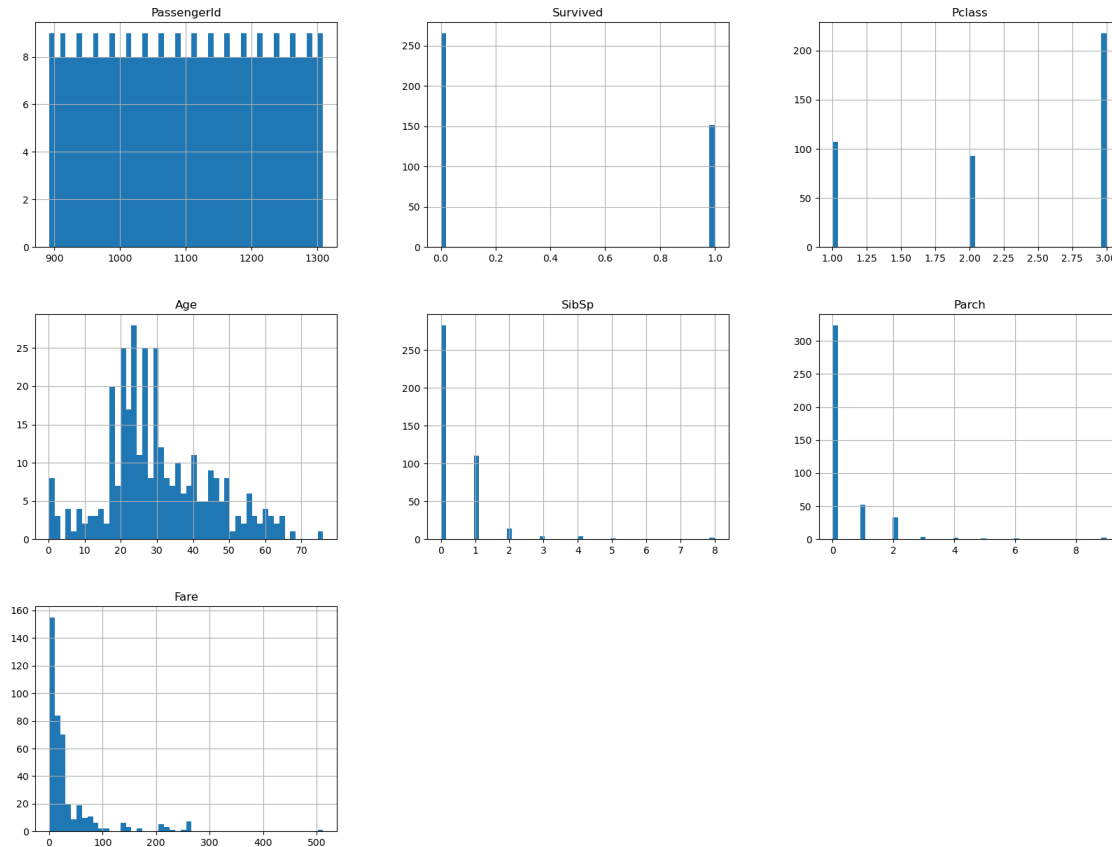
      Parch      Fare
count    418.000000  417.000000
mean         0.392344   35.627188
std         0.981429   55.907576
min          0.000000    0.000000
25%          0.000000    7.895800
50%          0.000000   14.454200
75%          0.000000   31.500000
max          9.000000  512.329200

```

```

[6]: titanic.hist(bins=50,figsize=(20,15))
plt.show()

```



```
[7]: train_set,test_set=train_test_split(titanic,random_state=100,test_size=0.
    ↪3,stratify=titanic['Survived'])

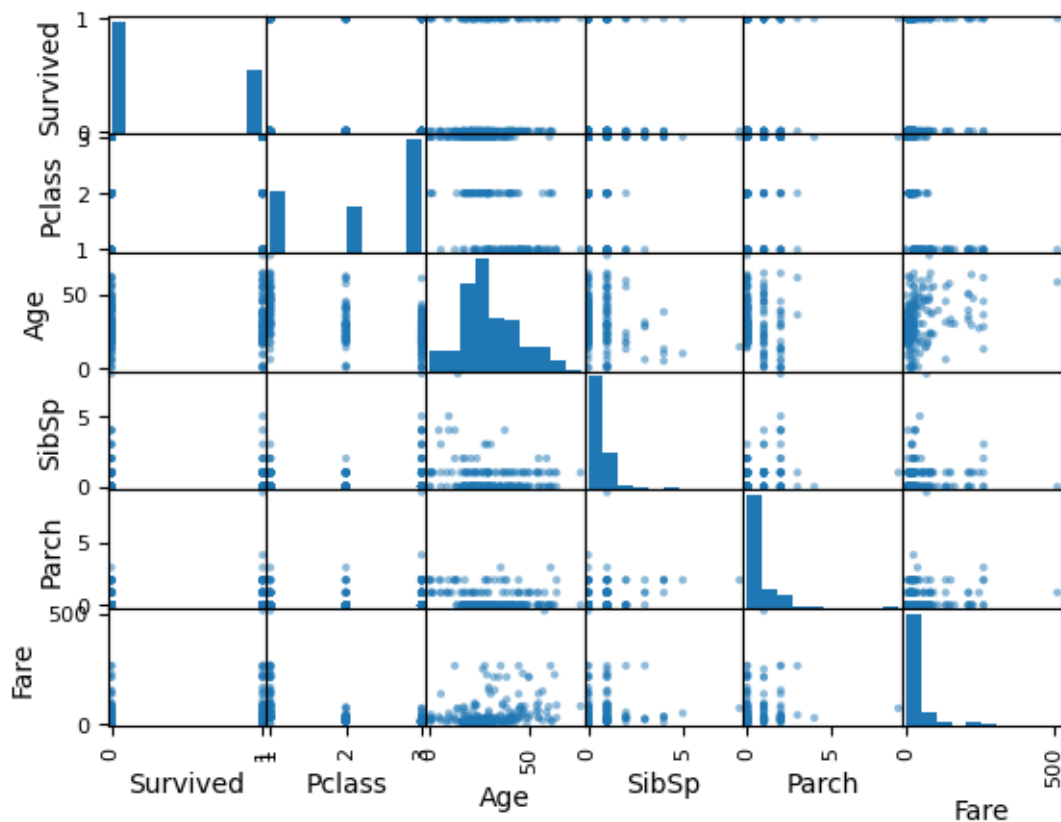
titanic=train_set.copy()
```

```
[8]: numeric_cols=['Survived','Pclass','Age','SibSp','Parch','Fare']
```

```
[9]: titanic[numeric_cols].corr()['Survived']
```

```
[9]: Survived    1.000000
Pclass        -0.108371
Age           0.054808
SibSp         0.052525
Parch         0.213548
Fare          0.204776
Name: Survived, dtype: float64
```

```
[10]: pd.plotting.scatter_matrix(titanic[numeric_cols])
plt.show()
```



```
[11]: labels,category=pd.factorize(titanic['Embarked'])
labels2,category2=pd.factorize(titanic['Sex'])
labels3,category3=pd.factorize(titanic['Cabin'])
```

```
[ ]:
```

```
[12]: titanic['embarked_num']=labels
titanic['Sex_num']=labels2
titanic['cabin_num']=labels3
```

```
[13]: numeric_cols.extend(['embarked_num','Sex_num','cabin_num'])
```

```
[14]: titanic[numeric_cols].corr()
```

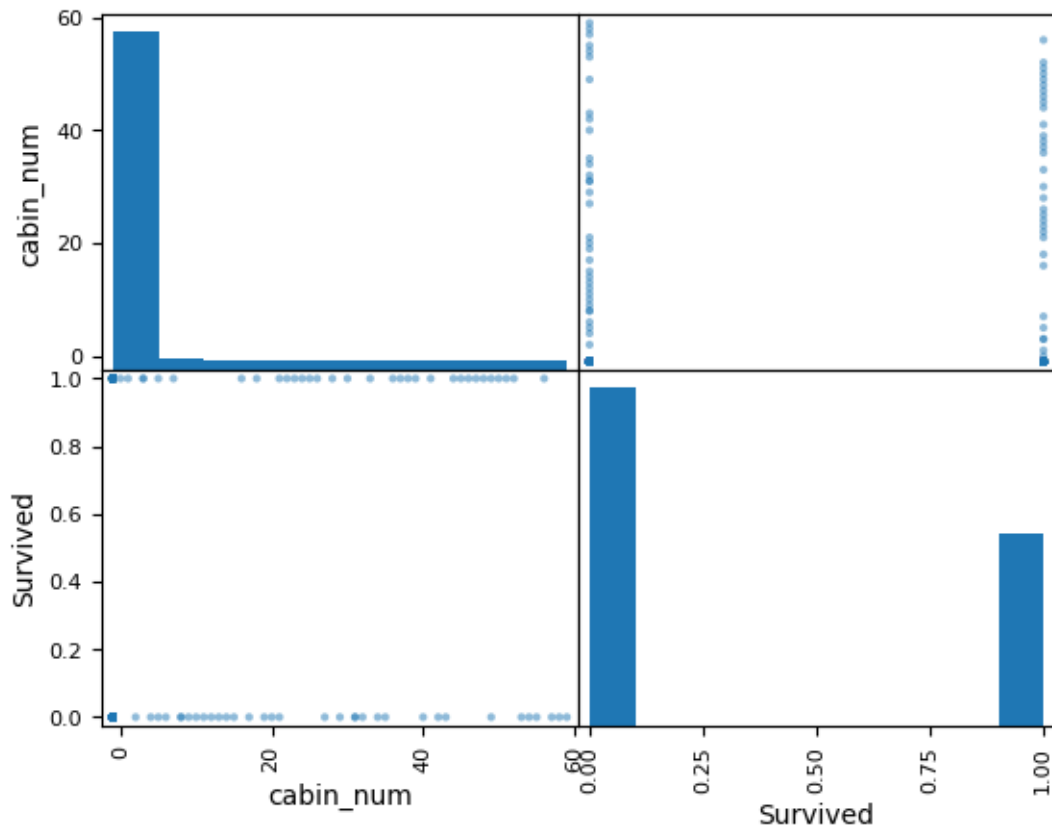
```
[14]:
```

	Survived	Pclass	Age	SibSp	Parch	Fare	\
Survived	1.000000	-0.108371	0.054808	0.052525	0.213548	0.204776	
Pclass	-0.108371	1.000000	-0.545526	0.028740	0.037316	-0.569527	
Age	0.054808	-0.545526	1.000000	-0.133935	-0.167283	0.308094	
SibSp	0.052525	0.028740	-0.133935	1.000000	0.356761	0.140274	
Parch	0.213548	0.037316	-0.167283	0.356761	1.000000	0.191723	

Fare	0.204776	-0.569527	0.308094	0.140274	0.191723	1.000000
embarked_num	0.139645	0.043861	0.131046	-0.106346	-0.108267	0.055062
Sex_num	1.000000	-0.108371	0.054808	0.052525	0.213548	0.204776
cabin_num	0.137130	-0.604828	0.390074	0.059850	0.086654	0.583569

	embarked_num	Sex_num	cabin_num
Survived	0.139645	1.000000	0.137130
Pclass	0.043861	-0.108371	-0.604828
Age	0.131046	0.054808	0.390074
SibSp	-0.106346	0.052525	0.059850
Parch	-0.108267	0.213548	0.086654
Fare	0.055062	0.204776	0.583569
embarked_num	1.000000	0.139645	0.027665
Sex_num	0.139645	1.000000	0.137130
cabin_num	0.027665	0.137130	1.000000

```
[15]: pd.plotting.scatter_matrix(titanic[['cabin_num', 'Survived']])
plt.show()
```



```
[16]: train_para=train_set.
      ↪drop(['PassengerId','Name','Ticket','Survived','Cabin'],axis=1) #
      ↪parameters for training
      train_label=train_set['Survived']           # labels for training
```

```
[ ]:
```

```
[17]: train_para.isnull().sum()
```

```
[17]: Pclass      0
      Sex       0
      Age      61
      SibSp     0
      Parch     0
      Fare      1
      Embarked  0
      dtype: int64
```

```
[18]: train_para.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 292 entries, 192 to 315
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Pclass      292 non-null    int64
1   Sex         292 non-null    object
2   Age         231 non-null    float64
3   SibSp       292 non-null    int64
4   Parch       292 non-null    int64
5   Fare        291 non-null    float64
6   Embarked    292 non-null    object
dtypes: float64(2), int64(3), object(2)
memory usage: 18.2+ KB
```

```
[19]: from sklearn.base import BaseEstimator, TransformerMixin
      class CleanNumericData(BaseEstimator, TransformerMixin):
          def __init__(self,method='median',fill_value=0):
              self.method=method
              self.fill_value=fill_value

          def give_numeric_cols(self,data):
              col_names=[data.columns[i] for i in range(len(data.dtypes)) if (data.
              ↪dtypes[i]==int or data.dtypes[i]==float)]
              return col_names

          def fit(self,data):
```

```

        return self

    def transform(self,data):
        imputer=SimpleImputer(strategy=self.method)
        num_cols=self.give_numeric_cols(data)
        X=imputer.fit_transform(data[num_cols])
        X_df=pd.DataFrame(X,columns=num_cols,index=data.index)
        return X_df

def ApplyOneHotEncoding(data):
    cato_cols=['Sex','Embarked']
    cat_encoder=OneHotEncoder()
    encoded=cat_encoder.fit_transform(data[cato_cols])
    encoded_df=pd.DataFrame(encoded.
    ↳toarray(),columns=['female','male','embarked_C','embarked_Q','embarked_S'],index=data.
    ↳index)
    encoded_df=encoded_df.drop('male',axis=1)
    return encoded_df

num_pipeline = Pipeline([
    ('cleaner',CleanNumericData()),
    ('standard_scaler',StandardScaler())
])

full_pipeline=ColumnTransformer([
    ('num',num_pipeline,['Pclass','Age','SibSp','Parch','Fare']),
    ('cat',OneHotEncoder(),['Sex','Embarked']),
])

```

```
[20]: temp=train_para.copy()
```

```
[21]: temp2=full_pipeline.fit_transform(temp)
```

```
[22]: transformed_train_para=pd.DataFrame(temp2,columns=['Pclass','Age','SibSp',
↳'Parch',
↳'Fare','Female','Male','Embarked_c','Embarked_S','Embarked_q'],index=temp.
↳index)
```

```
[23]: transformed_train_para
```

```
[23]:
```

	Pclass	Age	SibSp	Parch	Fare	Female	Male	\
192	0.890640	-1.444960	0.636853	0.748420	-0.384944	0.0	1.0	
134	0.890640	1.023488	-0.476686	-0.426519	-0.494953	0.0	1.0	
317	-0.275579	-0.857234	-0.476686	-0.426519	-0.451573	0.0	1.0	
135	0.890640	-0.465417	-0.476686	-0.426519	-0.495646	0.0	1.0	



```

72    0.890640 -0.073600 -0.476686 -0.426519 -0.494466    1.0    0.0
..      ...      ...      ...      ...      ...      ...
181 -1.441799  0.553308  0.636853  0.748420  0.758729    0.0    1.0
283  0.890640 -1.640869  0.636853  0.748420 -0.372520    1.0    0.0
390 -1.441799 -0.543780 -0.476686 -0.426519  0.930995    0.0    1.0
137 -0.275579 -0.308690 -0.476686 -0.426519 -0.409930    0.0    1.0
315  0.890640 -1.013961 -0.476686 -0.426519 -0.497660    1.0    0.0

```

```

      Embarked_c Embarked_S Embarked_q
192          0.0          0.0          1.0
134          0.0          0.0          1.0
317          0.0          0.0          1.0
135          0.0          0.0          1.0
72           0.0          0.0          1.0
..      ...      ...      ...
181          1.0          0.0          0.0
283          1.0          0.0          0.0
390          0.0          0.0          1.0
137          0.0          0.0          1.0
315          0.0          1.0          0.0

```

[292 rows x 10 columns]

```
[24]: test_param=test_set.drop('Survived',axis=1)
      test_labels=test_set['Survived']
```

```
[25]: temp=test_param.copy()
```

```
[26]: temp2=full_pipeline.fit_transform(temp)
```

```
[27]: transformed_test_para=pd.DataFrame(temp2,columns=['Pclass','Age', 'SibSp',
↳ 'Parch',
↳ 'Fare','Female','Male','Embarked_c','Embarked_S','Embarked_q'],index=temp.
↳ index)
```

```
[28]: transformed_test_para
```

```

[28]:      Pclass      Age      SibSp      Parch      Fare  Female  Male  \
131 -1.67332  1.900885 -0.553748 -0.375548 -0.053824    0.0    1.0
45   0.83666 -0.348615 -0.553748 -0.375548 -0.521055    0.0    1.0
348 -0.41833 -0.428954 -0.553748 -0.375548 -0.394454    0.0    1.0
118 -1.67332  0.535117 -0.553748 -0.375548  1.007620    0.0    1.0
173  0.83666 -0.268276 -0.553748 -0.375548 -0.536856    0.0    1.0
..      ...      ...      ...      ...      ...      ...
364 -1.67332 -0.348615  0.571611 -0.375548  0.557987    1.0    0.0
322 -0.41833 -0.268276 -0.553748 -0.375548 -0.405808    0.0    1.0
1    0.83666  1.418849  0.571611 -0.375548 -0.542061    1.0    0.0

```

```
120 -0.41833 -1.393026 -0.553748 -0.375548 -0.343360      1.0    0.0
406 -0.41833 -0.509294  0.571611 -0.375548 -0.462580      0.0    1.0
```

```
      Embarked_c Embarked_S Embarked_q
131          1.0          0.0          0.0
45           0.0          0.0          1.0
348          0.0          0.0          1.0
118          1.0          0.0          0.0
173          1.0          0.0          0.0
..          ...          ...          ...
364          1.0          0.0          0.0
322          0.0          0.0          1.0
1           0.0          0.0          1.0
120          0.0          0.0          1.0
406          0.0          0.0          1.0
```

[126 rows x 10 columns]

```
[29]: ds_tree=DecisionTreeClassifier(max_depth=1)
      kfold=KFold(n_splits=5,random_state=100,shuffle=True)
      score=cross_val_score(estimator=ds_tree,X=transformed_train_para,y=train_label,cv=kfold)
```

## 11.1 Random Forest

```
[30]: rnd_clf=RandomForestClassifier(n_estimators=500,max_leaf_nodes=16,n_jobs=-1)
```

```
[31]: rnd_clf.fit(transformed_train_para,train_label)
```

```
[31]: RandomForestClassifier(max_leaf_nodes=16, n_estimators=500, n_jobs=-1)
```

```
[32]: rnd_clf_pred=rnd_clf.predict(transformed_test_para)
```

## 11.2 Bagging

```
[33]: bag_clf=BaggingClassifier(DecisionTreeClassifier(splitter='random',max_leaf_nodes=16),n_estimators=500,bootstrap=True,n_jobs=-1)
```

```
[34]: bag_clf.fit(transformed_train_para,train_label)
```

```
[34]: BaggingClassifier(estimator=DecisionTreeClassifier(max_leaf_nodes=16,
                                                         splitter='random'),
                       n_estimators=500, n_jobs=-1)
```

```
[35]: bag_clf_pred=bag_clf.predict(transformed_test_para)
```

### 11.3 AdaBoost

```
[36]: ada_clf=AdaBoostClassifier(DecisionTreeClassifier(max_depth=1),n_estimators=200,algorithm='SAM
↳R',learning_rate=0.5)
```

```
[37]: ada_clf.fit(transformed_train_para,train_label)
```

```
[37]: AdaBoostClassifier(estimator=DecisionTreeClassifier(max_depth=1),
learning_rate=0.5, n_estimators=200)
```

```
[38]: ada_clf_pred=ada_clf.predict(transformed_test_para)
```

### 11.4 Gradient boosting

```
[39]: gb_clf=GradientBoostingClassifier(max_depth=2,n_estimators=3,learning_rate=1.0)
```

```
[40]: gb_clf.fit(transformed_train_para,train_label)
```

```
[40]: GradientBoostingClassifier(learning_rate=1.0, max_depth=2, n_estimators=3)
```

```
[41]: gb_clf_pred=gb_clf.predict(transformed_test_para)
```

### 11.5 XGBoost

```
[42]: xgb_model=xgb.XGBClassifier()
```

```
[43]: xgb_model.fit(transformed_train_para,train_label)
```

```
[43]: XGBClassifier(base_score=None, booster=None, callbacks=None,
colsample_bylevel=None, colsample_bynode=None,
colsample_bytree=None, device=None, early_stopping_rounds=None,
enable_categorical=False, eval_metric=None, feature_types=None,
gamma=None, grow_policy=None, importance_type=None,
interaction_constraints=None, learning_rate=None, max_bin=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=None, n_jobs=None,
num_parallel_tree=None, random_state=None, ...)
```

```
[44]: xgb_pred=xgb_model.predict(transformed_test_para)
xgb_pred
```

```
[44]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1,
1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1,
0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0,
```

```
0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0])
```

## 11.6 Blending

```
[45]: class blending_ensemble:
    def __init__(self, meta_model, models):
        self.meta_model = meta_model
        self.models = models
        self.data = None
        self.labels = None
        self.tst_size = None

    def get_meta_train_data(self):
        train_meta = []
        from sklearn.model_selection import train_test_split
        train_para, val_para, train_labels, val_labels = train_test_split(
            self.data, self.labels, test_size=self.tst_size, stratify=self.labels)
        for model in models:
            model.fit(train_para, train_labels)
            preds = model.predict(val_para)
            preds = preds.reshape(len(preds), 1)
            train_meta.append(preds)

        train_meta = np.hstack(train_meta)

        return (train_meta, val_labels)

    def get_meta_test_data(self, predict_data):
        test_meta = []
        for model in models:
            preds = model.predict(predict_data)
            preds = preds.reshape(len(preds), 1)
            test_meta.append(preds)
        test_meta = np.hstack(test_meta)
        return test_meta

    def fit(self, train_data, train_labels, test_size):
        self.data = train_data
        self.labels = train_labels
        self.tst_size = test_size

        meta_train_data, meta_train_labels = self.get_meta_train_data()
        # print(meta_train_data)
        self.meta_model.fit(meta_train_data, meta_train_labels)
```

```

def predict(self, predict_data):

    meta_test_data=self.get_meta_test_data(predict_data)
    # print(meta_test_data)
    return self.meta_model.predict(meta_test_data)

```

```

[46]: models=[DecisionTreeClassifier(),KNeighborsClassifier(),LogisticRegression(),SVC(probability=True)]

      bl_en=blending_ensemble(meta_model=LogisticRegression(),models=models)

```

```

[47]: bl_en.fit(transformed_train_para,train_label,test_size=0.3)

```

```

[48]: bl_pred=bl_en.predict(transformed_test_para)

```

## 11.7 Stacking

```

[49]: class stacking_ensemble:
      def __init__(self,models,meta_model):
          self.data=None
          self.models=models
          self.meta_model=meta_model
          self.labels=None
          self.number_of_folds=None

      def get_folds(self):
          start=0
          fold_size=len(self.data)//self.number_of_folds
          folds=[]
          for j in range(self.number_of_folds):
              if j==(self.number_of_folds-1):
                  # print(start,start+fold_size)
                  hold_out=self.data[start:]
                  hold_out_labels=self.labels[start:]

                  train_on=self.data[:start]
                  train_on_labels=self.labels[:start]

              else:

                  hold_out=self.data[start:(start+fold_size)]
                  hold_out_labels=self.labels[start:(start+fold_size)]

                  train_on=np.concatenate(( self.data[:start] , self.
→data[(start+fold_size):] ) ,axis=0)

```

```

        train_on_labels=np.concatenate(( self.labels[:start] , self.
↪labels[(start+fold_size):] ), axis=0)

        start+=fold_size

        folds.append((train_on,train_on_labels,hold_out,hold_out_labels))

    return folds

def get_meta_train_data(self):
    train_meta=[]

    folds=self.get_folds()
    train_meta_labels=np.array([])

#         getting labels for validation set, this will be used in
↪stacking_ensemble class
    for i in folds:
        train_meta_labels=np.concatenate((train_meta_labels,i[3]),axis=0)

#         multiple models
    if type(self.models)==list:

        for model in models:

            pred=np.array([])
            pred=pred.reshape(len(pred),1)

            for i in folds:
#                 train base model
                model.fit(i[0],i[1])
#                 predict for validation set
                curr_pred=model.predict(i[2])

                curr_pred=curr_pred.reshape(len(curr_pred),1)
                pred=np.concatenate((pred,curr_pred),axis=0)

            train_meta.append(pred)
            train_meta=np.hstack(train_meta)

#         single model
    else:
        print('y')
        pred=np.array([])
        pred=pred.reshape(len(pred),1)
        for i in folds:

```

```

        self.models.fit(i[0],i[1])
        curr_pred=self.models.predict(i[2])
        curr_pred=curr_pred.reshape(len(curr_pred),1)
        pred=np.concatenate((pred,curr_pred),axis=0)
        train_meta.append(pred)

    return (train_meta,train_meta_labels)

def get_meta_test_data(self,predict_data):
    test_meta=[]

    if type(self.models)==list:

        pred=np.array([])
        pred=pred.reshape(len(pred),1)

        for model in models:

            curr_pred=model.predict(predict_data)
            curr_pred=curr_pred.reshape(len(curr_pred),1)

            test_meta.append(curr_pred)

        test_meta=np.hstack(test_meta)

    # single model
    else:

        pred=np.array([])
        pred=pred.reshape(len(pred),1)
        curr_pred=models.predict(predict_data)
        curr_pred=curr_pred.reshape(len(curr_pred),1)

        test_meta.append(curr_pred)

    return test_meta

def fit(self,data,labels,number_of_folds=5):

    self.data=np.array(data)
    self.labels=np.array(labels)
    self.number_of_folds=number_of_folds

```

```

        meta_train_data,meta_train_labels=self.get_meta_train_data()
        # print(meta_train_data)
        self.meta_model.fit(meta_train_data,meta_train_labels)

    def predict(self,predict_data):
        # temp=kfold(data=predict_data, labels=labels,
        ↪models=base_layer_models, number_of_folds=number_of_folds)
        predict_data=np.array(predict_data)
        meta_test_data=self.get_meta_test_data(predict_data)
        # print(meta_test_data)
        return self.meta_model.predict(meta_test_data)

```

```
[50]: models=[DecisionTreeClassifier(),KNeighborsClassifier(),LogisticRegression(),SVC(probability=True)]
      stk=stacking_ensemble(meta_model=LogisticRegression(),models=models)
```

```
[51]: stk.fit(data=transformed_train_para,labels=train_label,number_of_folds=10)
```

```
[52]: stk_pred=stk.predict(transformed_test_para)
```

```
[53]: from sklearn.metrics import recall_score,accuracy_score,precision_score,f1_score
```

## 11.8 Accuracy

```
[54]: print(f'bagging: {accuracy_score(y_true=test_labels, y_pred=bag_clf_pred)}')
      print(f'gradient boosting: {accuracy_score(y_true=test_labels,
      ↪y_pred=gb_clf_pred)}')
      print(f'AdaBoosting: {accuracy_score(y_true=test_labels, y_pred=ada_clf_pred)}')
      print(f'xgboosting: {accuracy_score(y_true=test_labels, y_pred=xgb_pred)}')
      print(f'blending: {accuracy_score(y_true=test_labels, y_pred=bl_pred)}')
      print(f'stacking: {accuracy_score(y_true=test_labels, y_pred=stk_pred)}')
```

```

bagging: 1.0
gradient boosting: 1.0
AdaBoosting: 1.0
xgboosting: 1.0
blending: 1.0
stacking: 1.0

```

## 11.9 precision

```
[55]: print(f'bagging: {precision_score(y_true=test_labels, y_pred=bag_clf_pred)}')
      print(f'gradient boosting: {precision_score(y_true=test_labels,
      ↪y_pred=gb_clf_pred)}')
      print(f'AdaBoosting: {precision_score(y_true=test_labels,
      ↪y_pred=ada_clf_pred)}')
      print(f'xgboosting: {precision_score(y_true=test_labels, y_pred=xgb_pred)}')
```



```
print(f'blending: {precision_score(y_true=test_labels, y_pred=bl_pred)}')
print(f'stacking: {precision_score(y_true=test_labels, y_pred=stk_pred)}')
```

```
bagging: 1.0
gradient boosting: 1.0
AdaBoosting: 1.0
xgboosting: 1.0
blending: 1.0
stacking: 1.0
```

## 11.10 Recall

```
[56]: print(f'bagging: {recall_score(y_true=test_labels, y_pred=bag_clf_pred)}')
      print(f'gradient boosting: {recall_score(y_true=test_labels,
      ↪y_pred=gb_clf_pred)}')
      print(f'AdaBoosting: {recall_score(y_true=test_labels, y_pred=ada_clf_pred)}')
      print(f'xgboosting: {recall_score(y_true=test_labels, y_pred=xgb_pred)}')
      print(f'blending: {recall_score(y_true=test_labels, y_pred=bl_pred)}')
      print(f'stacking: {recall_score(y_true=test_labels, y_pred=stk_pred)}')
```

```
bagging: 1.0
gradient boosting: 1.0
AdaBoosting: 1.0
xgboosting: 1.0
blending: 1.0
stacking: 1.0
```

## 12 F1

```
[57]: print(f'bagging: {f1_score(y_true=test_labels, y_pred=bag_clf_pred)}')
      print(f'gradient boosting: {f1_score(y_true=test_labels, y_pred=gb_clf_pred)}')
      print(f'AdaBoosting: {f1_score(y_true=test_labels, y_pred=ada_clf_pred)}')
      print(f'xgboosting: {f1_score(y_true=test_labels, y_pred=xgb_pred)}')
      print(f'blending: {f1_score(y_true=test_labels, y_pred=bl_pred)}')
      print(f'stacking: {f1_score(y_true=test_labels, y_pred=stk_pred)}')
```

```
bagging: 1.0
gradient boosting: 1.0
AdaBoosting: 1.0
xgboosting: 1.0
blending: 1.0
stacking: 1.0
```

## 12.1 Kfold Cross Validation

```
[62]: scores=cross_val_score(rnd_clf,transformed_train_para,train_label,  
                             scoring='neg_mean_squared_error',cv=5)  
scores=np.sqrt(-scores)
```

```
[63]: scores
```

```
[63]: array([0., 0., 0., 0., 0.])
```

```
[ ]:
```

## 12.2 Hold-out validation

```
[69]: from sklearn.metrics import mean_squared_error
```

```
[65]: X_train, X_test, y_train,   
      ↪ y_test=train_test_split(transformed_train_para,train_label,test_size=0.3)
```

```
[67]: rnd_clf.fit(X_train,y_train)
```

```
[67]: RandomForestClassifier(max_leaf_nodes=16, n_estimators=500, n_jobs=-1)
```

```
[68]: pred=rnd_clf.predict(X_test)
```

```
[70]: mean_squared_error(y_true=y_test,y_pred=pred)
```

```
[70]: 0.0
```

```
[ ]:
```